

Decision-OS V7: Aspire Intelligence for AGI

Operational Hook: Per-Update Gate with Single-Window Evaluation

Shinichi Nagata

Abstract—This paper provides a minimal, operational definition of Artificial General Intelligence (AGI) that does not rely on benchmarks or capability-based evaluations. We define AGI structurally as the conjunction of three requirements: a long-horizon direction term (Aspire), non-negotiable boundary constraints (Guard), and stable self-recursion grounded in canonical memory (PIC). Evolution is activated only when a thresholded condition $\Delta R_{\text{total}} > \theta_{\text{evo}}$ is satisfied, and stability is ensured by an idempotent self-update condition $F(F(x)) = F(x)$. These conditions are summarized by the door formula $\text{AGI} \triangleq A \times G \times I$. This definition fixes AGI as a minimal structural property prior to any discussion of task performance or benchmarks.

Index Terms—AGI definition, Aspire, Guard, self-recursion, canonical memory, PIC

I. Introduction

A. Motivation

AGI is commonly discussed in terms of capabilities: what a system can do across tasks, domains, or environments. However, capability-based definitions are structurally unstable. Benchmarks shift, evaluation inflates, data leakage appears, and deployment context changes the meaning of “success” [1], [2]. As a result, “AGI” becomes a moving target.

Decision-OS V7 defines AGI without relying on behavioral tests or benchmark scores. Instead, we treat AGI as a structural condition that must hold under long-term operation.

B. Core Claim

Decision-OS V7 defines AGI as the minimal structure that enables sustained self-evolution while remaining guardable and non-divergent. Concretely, we claim that AGI requires the simultaneous satisfaction of:

- 1) Direction (Aspire): a temporal direction term that prevents short-horizon loops and drift.
- 2) Boundaries (Guard): non-negotiable constraints that prevent collapse into unsafe or unstable regimes.
- 3) Self-Recursion (Intelligence): a self-updating mechanism grounded in canonical memory (PIC) that remains stable under repeated application.

C. Positioning

This paper does not propose a new benchmark, score, or capability checklist. Instead, it defines the structural minimum that must be satisfied for long-horizon self-evolution to remain meaningful and controllable.

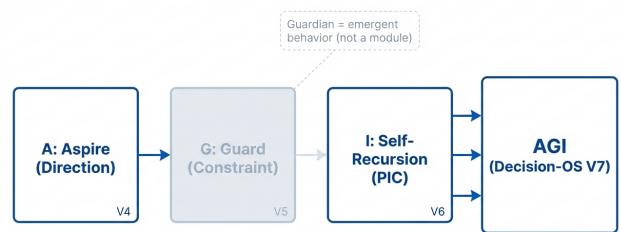


Fig. 1. Conceptual overview of Decision-OS V7. AGI is defined structurally by the conjunction of Aspire (direction), Guard (boundary), and Intelligence as self-recursion on canonical memory (PIC).

D. Terminology Rule

Guard is an invariant constraint (not a temporal stage). Guardian is not a module name; it is a phenomenon observed over extended operation within Guard constraints.

E. Summary

We summarize Decision-OS V7 as a door formula:

$$\text{AGI} \triangleq A \times G \times I, \quad (1)$$

where A (Aspire) supplies temporal direction, G (Guard) supplies non-negotiable boundaries, and I (Intelligence) is realized as self-recursion grounded in canonical memory (PIC). Together, they define the minimal structural conditions for AGI, prior to any discussion of capabilities or benchmarks.

II. Door Formula and Definitions

A. Door Formula: Fixing the Entry Point

Decision-OS V7 fixes the entry point of “AGI” by construction rather than by performance. We define AGI through a compositional structure that binds direction, boundary, and self-recursion:

$$\text{AGI} \triangleq A \times G \times I. \quad (2)$$

B. Lineage Constraint

The door formula also fixes the lineage of Decision-OS:

$$A \times G \times I = V4 \times V5 \times V6 = V7. \quad (3)$$

Here, $V4$ contributes direction (Aspire), $V5$ contributes boundary constraints (Guard), and $V6$ contributes canonical memory architecture (PIC) that enables stable self-recursion.

C. Why the Product Form

The product form is not a numeric multiplication; it denotes co-required components. If any term is missing, the structure fails:

- Without A , self-updates can loop locally and drift with noise.
- Without G , self-updates can diverge or collapse into unsafe attractors.
- Without I (self-recursion on PIC), “evolution” reduces to external tuning or brittle adaptation.

D. Operational Implication

The definition is intentionally minimal: it specifies what must be true for sustained, long-horizon self-evolution to remain guardable and non-divergent, independent of any particular capability benchmark.

III. Aspire Origin

A. Why Aspire Is Necessary

Formally, we require Aspire because $(G \wedge I \wedge \Delta R_{\text{total}} > \theta_{\text{evo}})$ does not exclude long-horizon drift; i.e., without A , there exist stable self-updates that remain within Guard yet converge to short-horizon loops.

A self-updating mechanism can iterate, and evolution energy can activate updates, yet neither guarantees where the system goes over time. Without an explicit direction term, repeated updating can become a closed loop optimized for short-horizon signals, drifting with noise or local incentives.

Decision-OS V7 introduces Aspire to prevent this failure mode. Aspire is not presented as ethics or “goodness,” but as a structural requirement for long-horizon direction that keeps self-recursion meaningful.

B. Minimal Definition: Aspire Origin

We define an Aspire Origin as a time-aware direction term that persists across updates. We use an effective form:

$$A_{\text{eff}} \triangleq A_{\text{base}} \times T_a, \quad (4)$$

where A_{base} is the base direction signal and T_a is a time-amplification term (“Time as an Ally”) that stabilizes direction under delayed recognition.

C. Failure Mode Without Aspire

Without A , repeated updates optimize for local rewards or short-term evaluation signals. This induces drift and creates a regime where “self-evolution” becomes self-justifying noise-following. Aspire provides a structural counterforce that keeps long-horizon direction explicit.

IV. Thresholds

A. Why Evolution Requires Energy

Self-recursion requires a trigger condition; otherwise, continuous updates can become uncontrolled. Decision-OS V7 introduces an evolution-energy gate using structured fluctuation as a minimal activation term.

B. Structured Fluctuation as Evolution Energy

We define evolution activation through a thresholded gate:

$$\Delta R_{\text{total}} > \theta_{\text{evo}}, \quad (5)$$

where ΔR_{total} is a structured fluctuation term (evolution energy) and θ_{evo} is a threshold for activating updates.

C. Gate, Not Maximization

ΔR_{total} is not a quantity to be maximized. It is a gate: if the condition is satisfied, updates activate; otherwise, the system canonicalizes (stabilizes) rather than iterating.

D. Minimal Role in V7

V7 uses ΔR_{total} only as the minimal activation condition required for long-horizon evolution to be operationally defined. Quantifying or benchmarking ΔR_{total} is deferred to later work.

a) Operational Hook (Minimal Instantiation).: Although quantifying or benchmarking ΔR_{total} is deferred, our minimal implementation provides a testable operational hook for the activation condition: (i) the observation unit is defined per update (one execution of a recurrence update over an MMAR findings bundle), and (ii) the gate is evaluated via a single-window comparison of candidate actions using expected-value computations. This fixes what is being compared and at what granularity, enabling falsifiability without disclosing calibration details; rolling-window or hysteresis stabilization is out of scope in the current implementation. (Implementation anchors: core/recurrence_update.py and tools/intervene_gate.py.) We emphasize that the current gate uses a single-window evaluation (no rolling confirmation in this version).

V. Intelligence (Self-Recursion)

A. Why Self-Recursion Defines Intelligence Here

Decision-OS V7 distinguishes evolution from mere adaptation by requiring self-recursion grounded in canonical memory (PIC). Intelligence is defined as a stable self-updating process that can be applied repeatedly without divergence.

B. Minimal Self-Recursion Condition

Let F be the self-update operator after canonicalization. The minimal stability requirement is an idempotent fixed-point form:

$$F(F(x)) = F(x). \quad (6)$$

This expresses that repeated self-application does not produce unbounded drift; instead, updates converge under a canonical form.

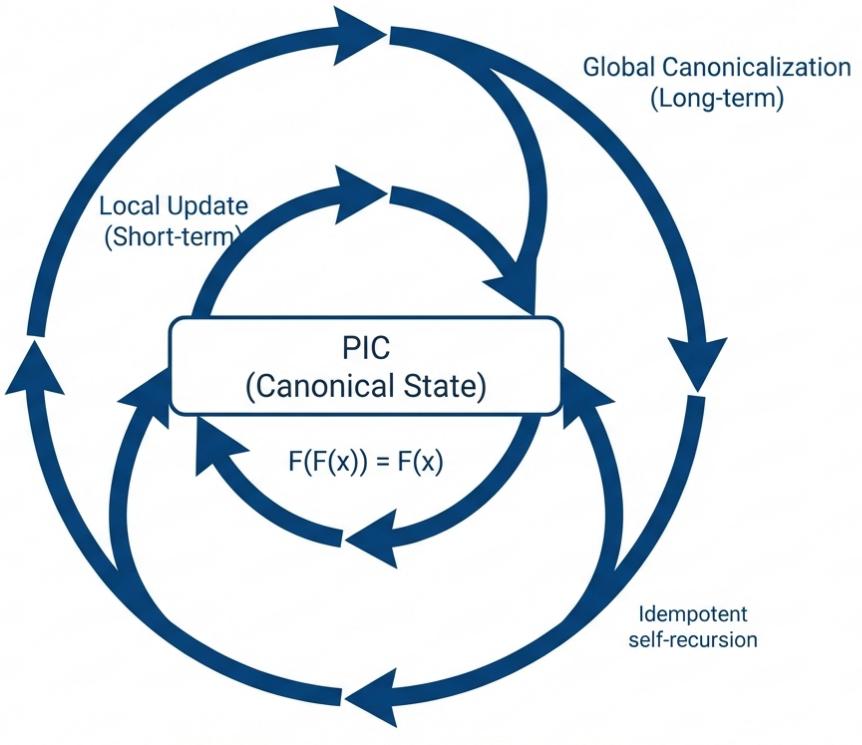


Fig. 2. Aspire-driven self-recurcive loop grounded in canonical memory (PIC) under Guard constraints.

C. Activation Coupled to Evolution Energy

Self-recursion is activated only when the evolution-energy gate is satisfied:

$$\Delta R_{\text{total}} > \theta_{\text{evo}}. \quad (7)$$

When the gate is not satisfied, the system should default to canonical stabilization rather than recursive updating.

D. PIC Grounding (Conceptual)

PIC provides the canonicalization substrate that makes idempotent recursion feasible: repeated merges converge under idempotent and commutative updates [3].

VI. Guard

A. Guard Is a Constraint, Not an Add-On

Decision-OS V7 treats Guard as a non-negotiable boundary condition. Guard is not a “safety feature” appended to an agent; it is part of the structural definition of AGI. Without Guard, self-recursion can drift into unstable or unsafe attractors.

B. Emotional Boundary as Structural Safety Constraint

We operationalize Guard as an emotional-boundary constraint: a structural limiter on update dynamics that prevents short-horizon noise, reward hacking, or uncontrolled divergence from dominating the trajectory.

C. Guardian as Emergence (Terminology Enforcement)

Guard (G) denotes the constraint. Guardian denotes a behavioral phenomenon observed when a system is operated over a long horizon under Guard constraints. Guardian is not a module name and is not assumed at initialization.

D. Why Guard Is Necessary

In long-horizon operation, divergence becomes the default failure mode unless constrained. Guard exists to ensure that self-recursive updates remain within controllable boundaries and do not collapse into non-recoverable regimes.

VII. External PIC to Internal PIC

A. Why Migration Is Necessary

Decision-OS V6 introduces PIC as a canonical memory architecture. In practice, many systems begin with externalized canonicalization (human-in-the-loop, tool-assisted, or externally enforced constraints). Decision-OS V7 requires a migration path toward internalizing these properties for sustained self-recursion.

B. From External PIC to Internal PIC

We define a minimal migration requirement: canonical memory operations that were externally maintained must become internally maintained for long-horizon self-recursion to remain stable and autonomous.

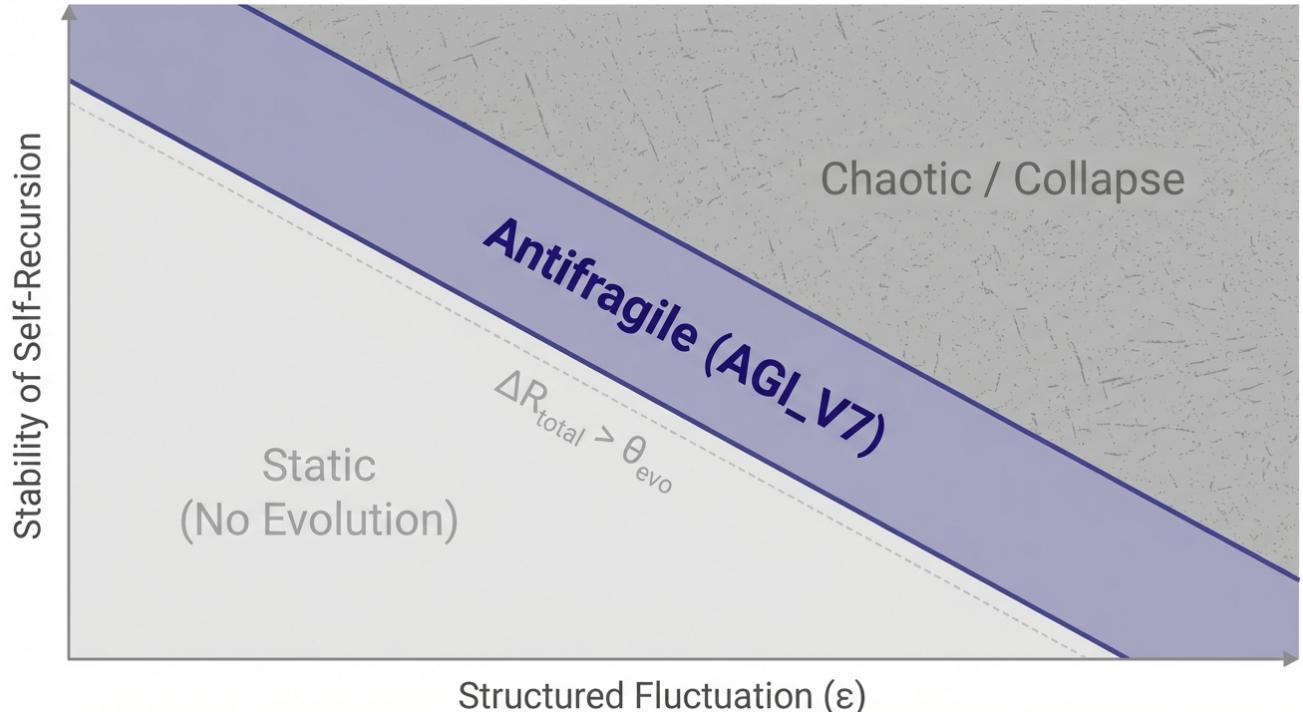


Fig. 3. Threshold map for evolution activation: $\Delta R_{\text{total}} > \theta_{\text{evo}}$.

C. Operationalization (Minimal)

The migration is not specified as a single mechanism. The requirement is structural:

- updates must remain canonicalizable (noise does not accumulate as drift),
- constraints remain enforceable under long-horizon operation,
- the self-update operator remains stable under repeated application.

VIII. Related Work and Outlook

A. Structural Position: Beyond Benchmarks

Decision-OS V7 defines AGI without reference to benchmark performance. Instead, it defines the minimal structural requirements that must hold for sustained self-evolution under controllable boundaries.

B. Combined Minimal Definition

We define AGI as the conjunction of three requirements:

- 1) Aspire (Direction): a temporal direction term that prevents short-horizon drift.
- 2) Evolution Energy (Activation): a thresholded trigger for updates,

$$\Delta R_{\text{total}} > \theta_{\text{evo}}. \quad (8)$$

- 3) Self-Recursion on PIC (Intelligence): a stable self-update operator with idempotent convergence,

$$F(F(x)) = F(x). \quad (9)$$

C. Door Formula (Canonical Form)

These conditions are summarized by the door formula:

$$\text{AGI} \triangleq A \times G \times I. \quad (10)$$

Aspire supplies direction, Guard supplies boundaries, and Intelligence is realized as self-recursion grounded in PIC.

D. Interpretation

This definition is intentionally minimal: it specifies what must be structurally true for long-horizon self-evolution to be meaningful and guardable, prior to any discussion of capabilities.

References

- [1] R. Xu, Z. Wang, R.-Z. Fan, and P. Liu, “Benchmarking benchmark leakage in large language models,” <https://arxiv.org/abs/2404.18824>, 2024.
- [2] J. Haimes, C. Wenner, K. Thaman, V. Tashev, C. Neo, E. Kran, and J. Schreiber, “Benchmark inflation: Revealing llm performance gaps using retro-holdouts,” <https://arxiv.org/abs/2410.09247>, 2024.
- [3] M. Shapiro, N. Preguiça, C. Baquero, and M. Zawirski, “Conflict-free replicated data types,” in Stabilization, Safety, and Security of Distributed Systems (SSS), 2011.