Decision-OS Note: Time-is-an-Ally v2

Value amplification with disciplined re-evaluation (T_a) and discoverability (D_a)

Minimal Definition (v2)
Time-is-an-Ally is not a forecasting formula. Time amplifies *re-evaluation* (not discovery) under external signals. We separate "time amplification" from "discoverability":

$$\text{Value}_{\text{later}} = T_a(t) \times D_a.$$

$T_a(t)$ increases only via re-evaluation under external signals; D_a does not increase by time alone.

Non-Use (misread prevention)
Not meme/price prediction. Not hindsight absolution. Not automatic amplification by time alone.

Two-Universe Model (salvage vs amplification)
Universe A: Long-tail salvage. AI may collect and store ideas as knowledge-graph nodes (higher pickup, not guaranteed impact).
Universe B: Impact amplification. Ideas become reusable structures and are recommended/deployed; **here** Time-is-an-Ally matters most.
Transition A→B occurs when: (i) structure is extracted, (ii) reuse is demonstrated, and (iii) external adoption begins.

Prototype Form (Genesis×Time)
Let Talent, Market, Influence, Genesis G_{index}, $t_{\text{as-of}}$ and $t_{\text{re-eval}}$, and k **(calibration constant)**: estimated from empirical re-evaluation patterns (domain-dependent), **not tuned post-hoc** to fit desired outcomes.

$$A_{\text{total}} = (\text{Talent} \times \text{Market} \times \text{Influence}) \times G_{\text{index}} \times \left[1 + k \cdot \log\left(\frac{t_{\text{re-eval}}}{t_{\text{as-of}}}\right)\right]$$

$$T_a(t) = 1 + k \cdot \log\left(\frac{t_{\text{re-eval}}}{t_{\text{as-of}}}\right), \qquad A_{\text{eff}} = A_{\text{base}} \times T_a(t).$$

Re-evaluation is allowed only under external signals.

D_a: Minimal 3-item implementation (raise discoverability/legibility)
One-sentence definition (EN) → One diagram (I/O arrows) → One falsifiable slice (minimal test).
D_a rises when the unit becomes reusable, not when the text becomes longer.

Capability-Ceiling-Aware Review (Reverse Ta)
Past decisions should be reviewed under the capability ceiling available at the time (model/tool limits), not judged by today's stronger models; updates must be expressed as forward-only deltas.

CAP v1 (anti goalpost-moving)
Claim: — Dependencies (max3): — Alert-if (max3): — Recheck-when: — Hedge (1): —

DFR v1 (Delta-first reply)
As-of: — Δ: — Residue: — Next step: —

Repository: decision-os-paper (GitHub) Zenodo (v1 DOI): 10.5281/zenodo.18158755